CS 188: Artificial Intelligence
Fall 2011

Advanced Applications:
Robotics / Vision / Language

Dan Klein – UC Berkeley

Many slides from Sebastian Thrun, Pieter Abbeel, Jitendra Malik

1

Announcements

- This week:
  - No sections this week
  - Office hours modified (moved earlier), see Piazza
  - Final contest cutoff is 8pm Wednesday
- Grades:
  - W1-3, P1-4, Midterm in glookup, please check
  - P5: full credit in grade computations

2

So Far: Foundational Methods

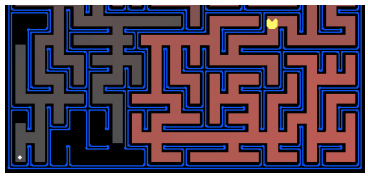

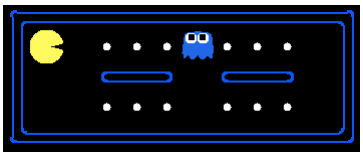

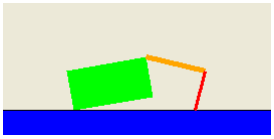

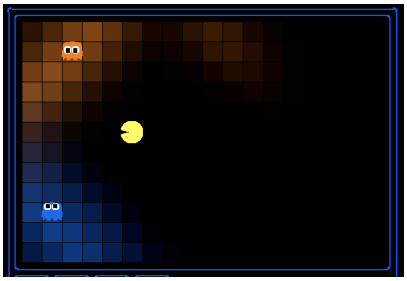

3

Now: Advanced Applications

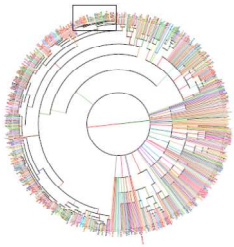

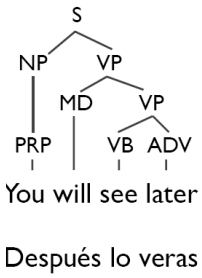

4

[DEMO: Race, Short]

Autonomous Vehicles

Autonomous vehicle slides adapted from Sebastian Thrun

[DEMO: GC Bad, Good]

Grand Challenge: Barstow, CA, to Primm, NV

- 150 mile off-road robot race across the Mojave desert
- Natural and manmade hazards
- No driver, no remote control
- No dynamic passing

An Autonomous Car

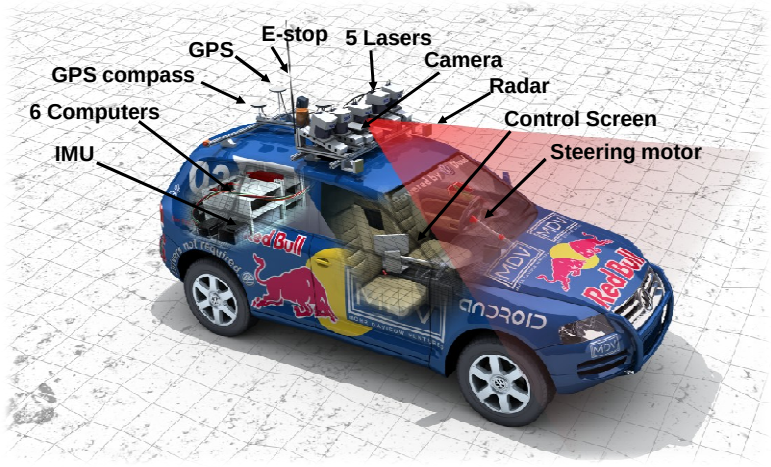

Actions: Steering Control

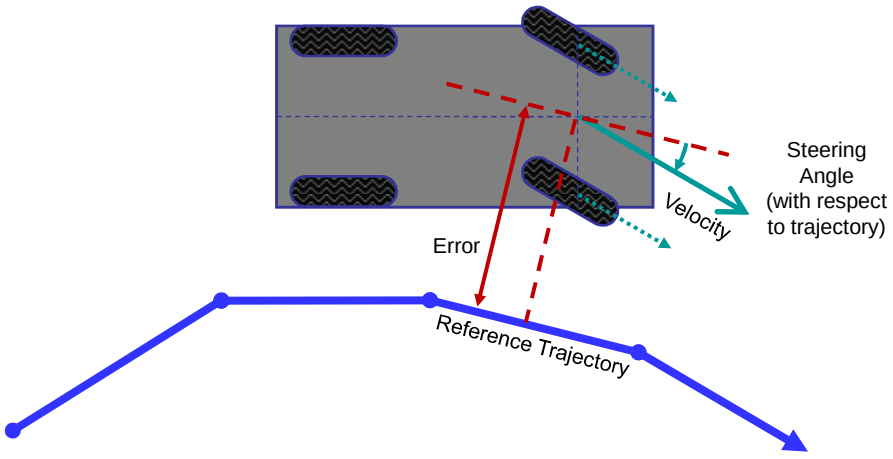

[DEMO: LIDAR]

Sensors: Laser Readings

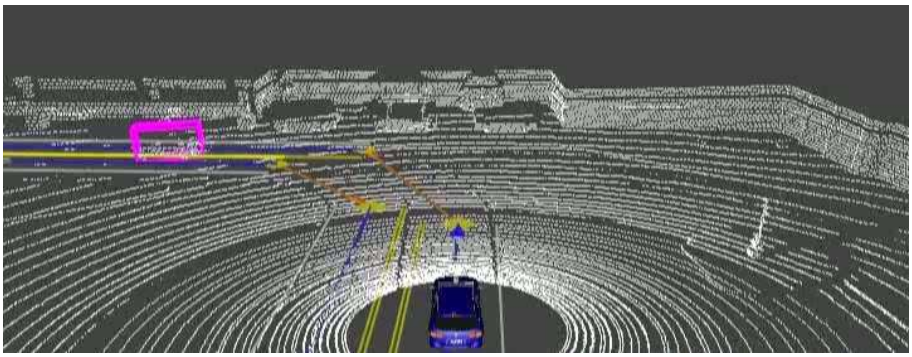

Readings: No Obstacles

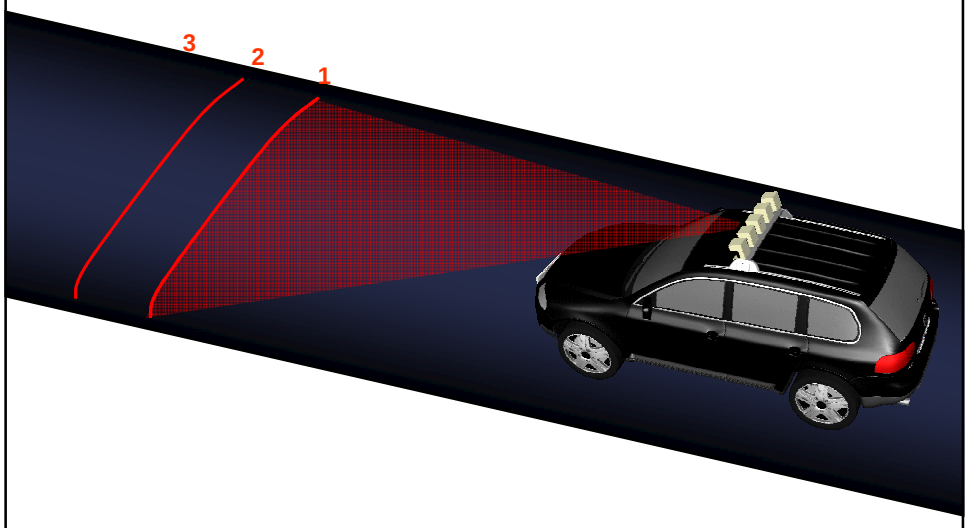

Readings: Obstacles

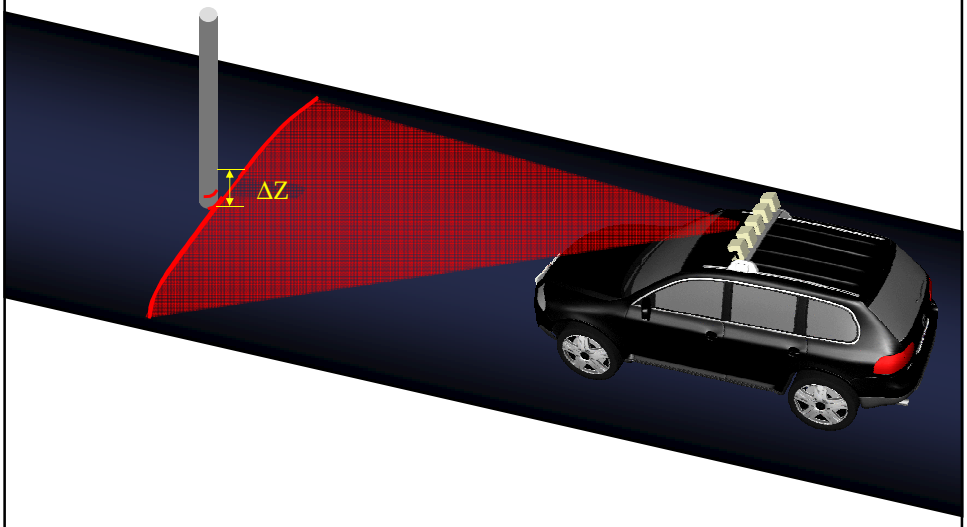

Obstacle Detection

Trigger if $|Z^i - Z^j| > 15\text{cm}$ for nearby z^i, z^j

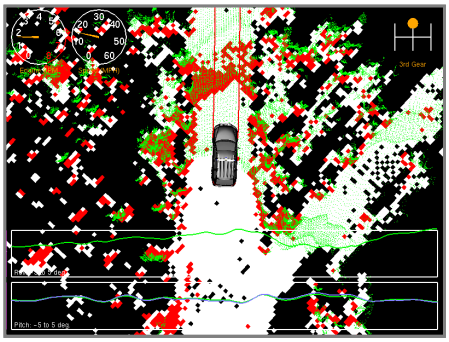

Raw Measurements: 12.6% false positives

Probabilistic Error Model

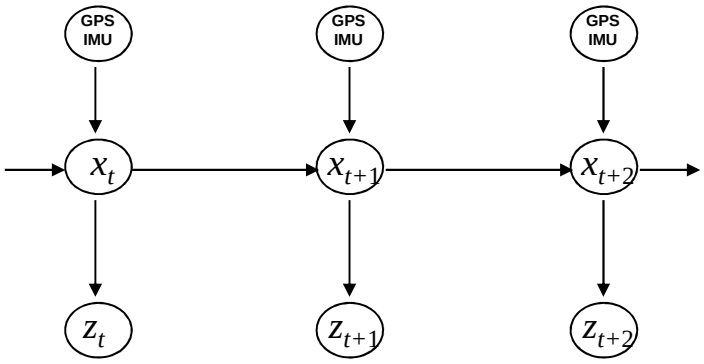

HMMs for Detection

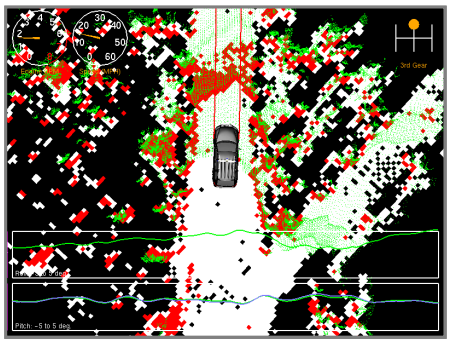

Raw Measurements: 12.6% false positives

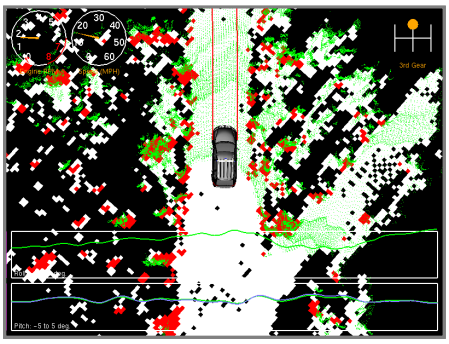

HMM Inference: 0.02% false positives

Environmental Tracking

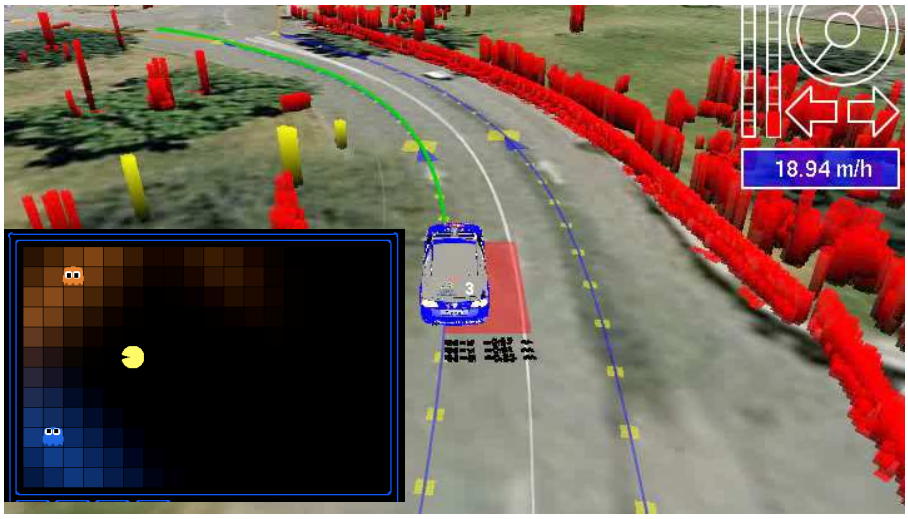

[DEMO: PEOPLE]

Sensors: Camera

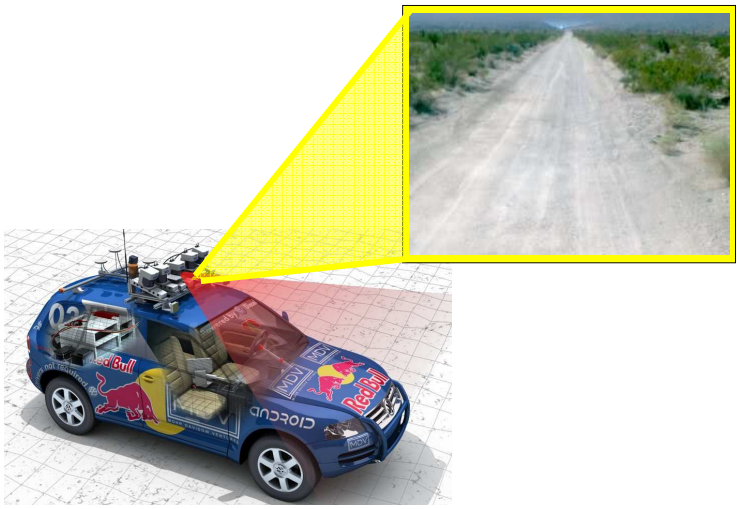

Object Recognition

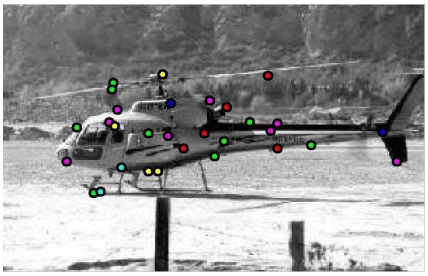

Template

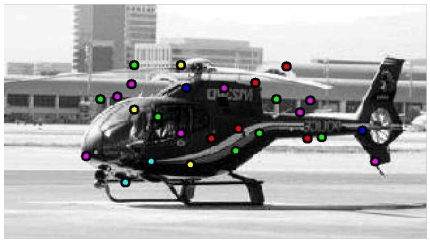

Query

Vision slides adapted from Jitendra Malik

Shape Context

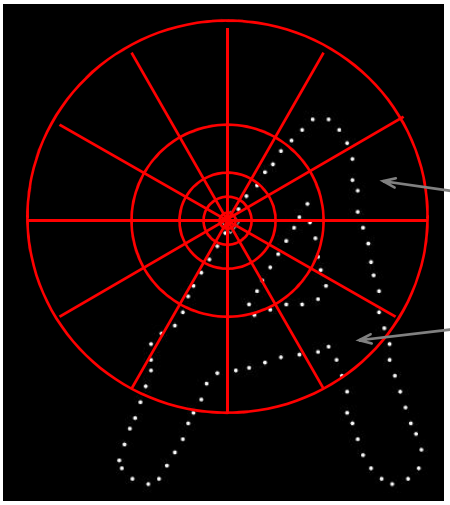

Count the number of points inside each bin, e.g.:

Count = 4

⋮

Count = 10

- Compact representation of distribution of points relative to each point

18

Shape Context

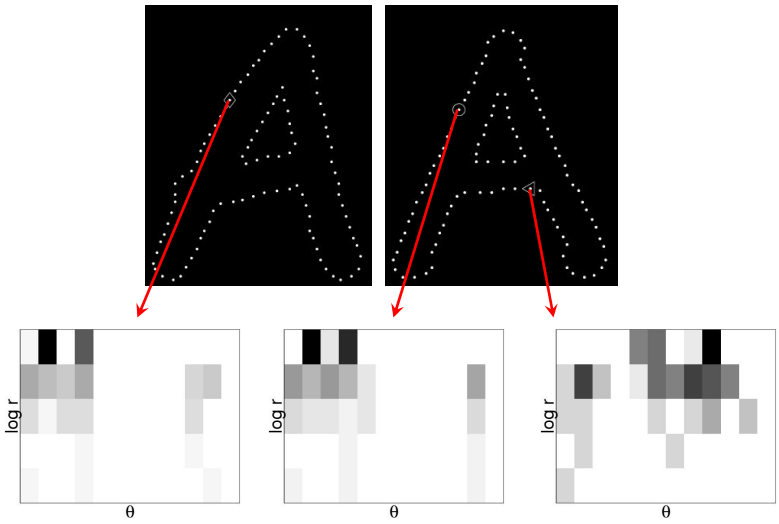

19

Similar Regions

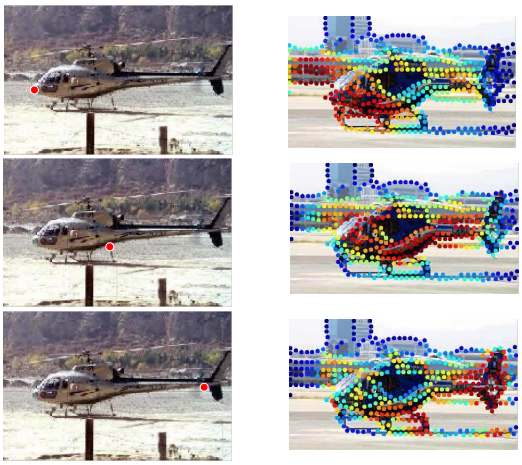

Color indicates similarity using local descriptors

20

Match for Image Similarity

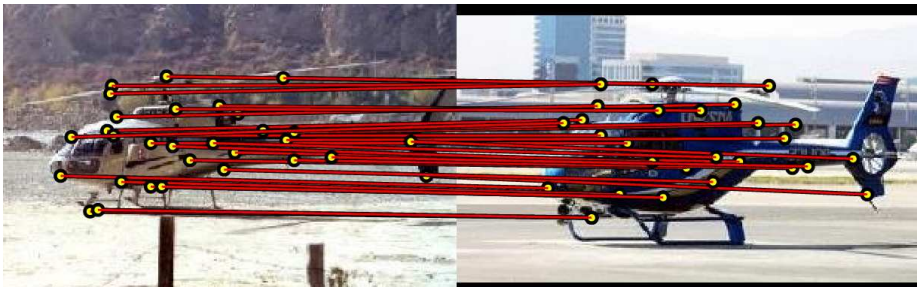

21

[DEMO: LIDAR 1]

Vision for a Car

[DEMO: LIDAR 2]

Self-Supervised Vision

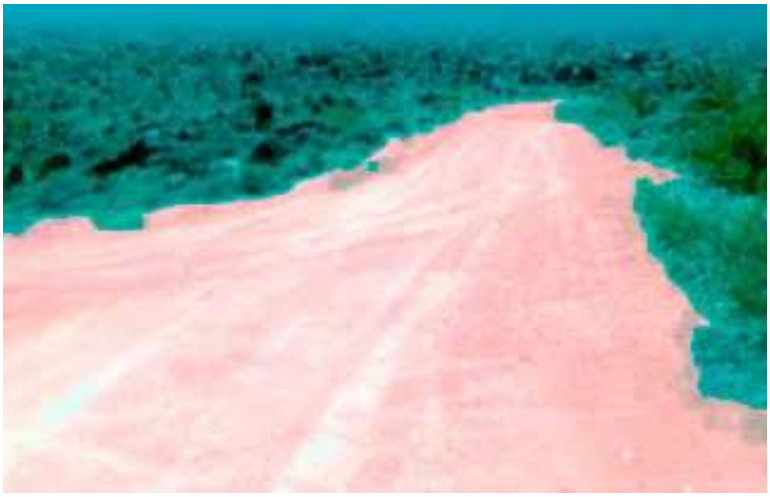

Complex Robot Control

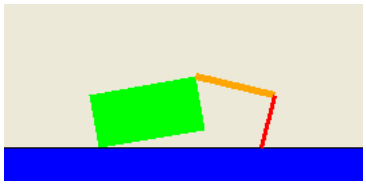

[demo – quad initial]

Robotic Control Tasks

- Perception / Tracking
  - Where exactly am I?
  - What's around me?
- Low-Level Control
  - How to move from position A to position B
  - Safety vs efficiency
- High-Level Control
  - What are my goals?
  - What are the optimal high-level actions?

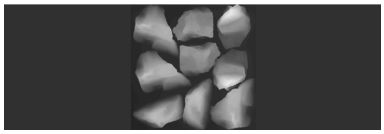

Low-Level Planning

- Low-level: move from configuration A to configuration B

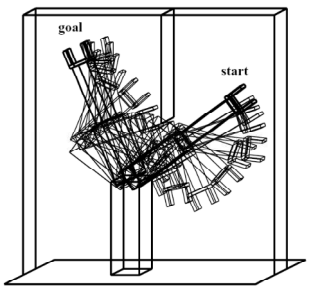

A Simple Robot Arm

- Configuration Space
  - What are the natural coordinates for specifying the robot's configuration?
  - These are the *configuration space* coordinates
  - Can't necessarily control all degrees of freedom directly
- Work Space
  - What are the natural coordinates for specifying the effector tip's position?
  - These are the *work space* coordinates

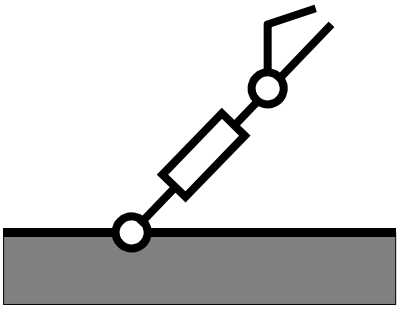

Coordinate Systems

- Workspace:
  - The world's (x, y) system
  - Obstacles specified here
- Configuration space
  - The robot's state
  - Planning happens here
  - Obstacles can be projected to here

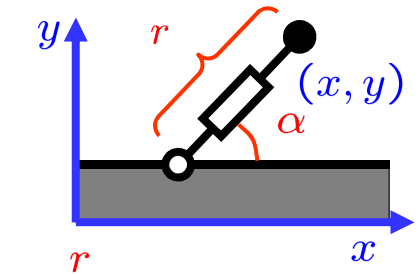

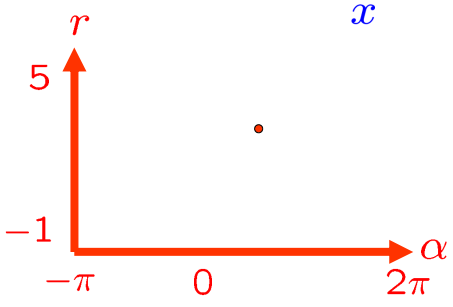

Obstacles in C-Space

- What / where are the obstacles?
- Remaining space is *free space*

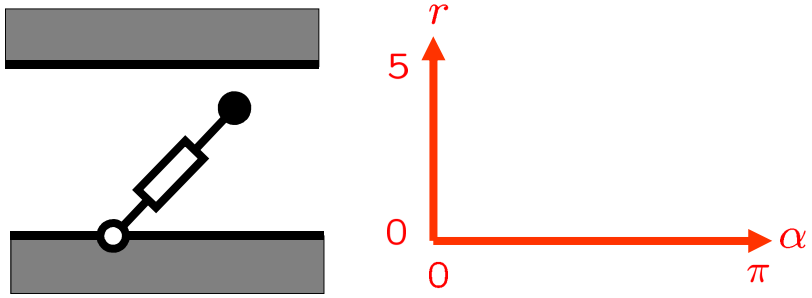

Example: A Less Simple Arm

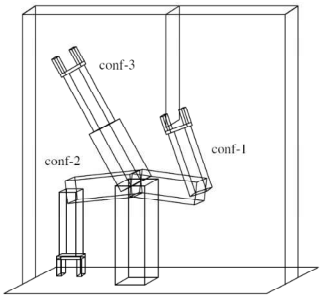

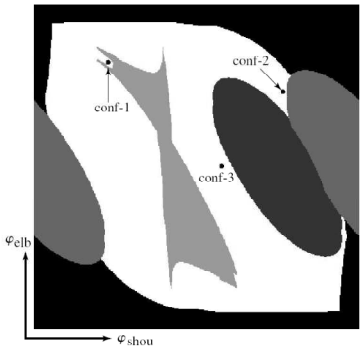

[DEMO]

Probabilistic Roadmaps

- Idea: sample random points as nodes in a visibility graph

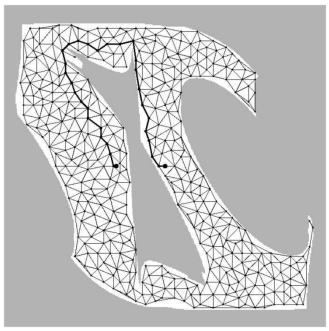

- This gives *probabilistic roadmaps*
  - Very successful in practice
  - Lets you add points where you need them
  - If insufficient points, incomplete or weird paths

High-Level Control

- Demonstrate path across the “training terrain”

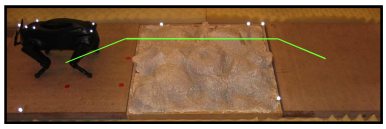

- Run apprenticeship learning to find a set of weights w
- Receive “testing terrain” (a height map)

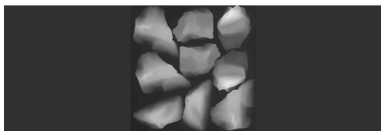

- Find a policy for crossing the testing terrain.

High DOF Robots [DEMOS]

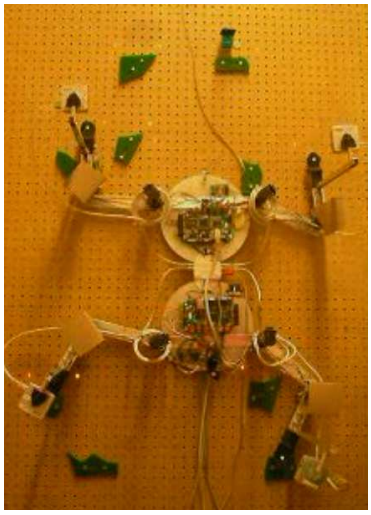

Videos from Pieter Abbeel, Jean-Claude Latombe